Analysis Report

void cellHistogramsNaive<unsigned char, unsigned char, int=64, int=8, int=8>(unsigned char*, unsigned char*, int, int, int)

|	
Duration	38.288 μs
Grid Size	[3,3,1]
Block Size	[16,16,1]
Registers/Thread	17
Shared Memory/Block	0 B
Shared Memory Requested	96 KiB
Shared Memory Executed	96 KiB
Shared Memory Bank Size	4 B

[0] GeForce GTX 960	
GPU UUID	GPU-0db32734-f94e-48a7-8b5d-4604317dc554
Compute Capability	5.2
Max. Threads per Block	1024
Max. Shared Memory per Block	48 KiB
Max. Registers per Block	65536
Max. Grid Dimensions	[2147483647, 65535, 65535]
Max. Block Dimensions	[1024, 1024, 64]
Max. Warps per Multiprocessor	64
Max. Blocks per Multiprocessor	32
Single Precision FLOP/s	2.644 TeraFLOP/s
Double Precision FLOP/s	82.624 GigaFLOP/s
Number of Multiprocessors	8
Multiprocessor Clock Rate	1.291 GHz
Concurrent Kernel	true
Max IPC	6
Threads per Warp	32
Global Memory Bandwidth	112.16 GB/s
Global Memory Size	4 GiB
Constant Memory Size	64 KiB
L2 Cache Size	1 MiB
Memcpy Engines	2
PCIe Generation	2
PCIe Link Rate	5 Gbit/s
PCIe Link Width	16

1. Compute, Bandwidth, or Latency Bound

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results below indicate that the performance of kernel "void cellHistogramsNaive<un..." is most likely limited by instruction and memory latency. You should first examine the information in the "Instruction And Memory Latency" section to determine how it is limiting performance.

1.1. Kernel Performance Is Bound By Instruction And Memory Latency

This kernel exhibits low compute throughput and memory bandwidth utilization relative to the peak performance of "GeForce GTX 960". These utilization levels indicate that the performance of the kernel is most likely limited by the latency of arithmetic or memory operations. Achieved compute throughput and/or memory bandwidth below 60% of peak typically indicates latency issues.

2. Instruction and Memory Latency

Instruction and memory latency limit the performance of a kernel when the GPU does not have enough work to keep busy. The results below indicate that the GPU does not have enough work because the kernel does not execute enough blocks.
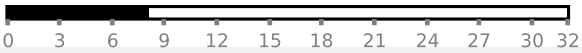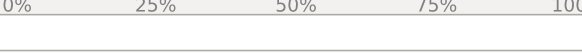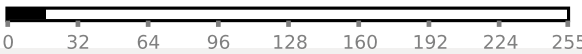
2.1. Grid Size Too Small To Hide Compute And Memory Latency

The kernel does not execute enough blocks to hide memory and operation latency. Typically the kernel grid size must be large enough to fill the GPU with multiple "waves" of blocks. Based on theoretical occupancy, device "GeForce GTX 960" can simultaneously execute 8 blocks on each of the 8 SMs, so the kernel may need to execute a multiple of 64 blocks to hide the compute and memory latency. If the kernel is executing concurrently with other kernels then fewer blocks will be required because the kernel is sharing the SMs with those kernels.

Optimization: Increase the number of blocks executed by the kernel.
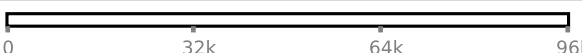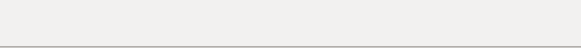
2.2. Occupancy Is Not Limiting Kernel Performance

The kernel's block size, register usage, and shared memory usage allow it to fully utilize all warps on the GPU.

Variable	Achieved	Theoretical	Device Limit	Grid Size: [3,3,1] (9 blocks) Block Size: [16,16,1] (256 threads)
Occupancy Per SM				
Active Blocks		8	32	
Active Warps	6.54	64	64	
Active Threads		2048	2048	
Occupancy	10.2%	100%	100%	
Warps				
Threads/Block		256	1024	
Warps/Block		8	32	
Block Limit		8	32	
Registers				
Registers/Thread		17	255	
Registers/Block		6144	65536	
Block Limit		10	32	
Shared Memory				
Shared Memory/Block		0	98304	
Block Limit			32	

2.3. Occupancy Charts

The following charts show how varying different components of the kernel will impact theoretical occupancy.

Varying Block Size

Warps Per SM vs **Threads Per Block**

64 @ 256

Varying Register Count

Warps Per SM vs **Registers Per Thread**

64 @ 17

4

Varying Shared Memory Usage

64 @ 0k

Shared Memory Per Block (bytes)
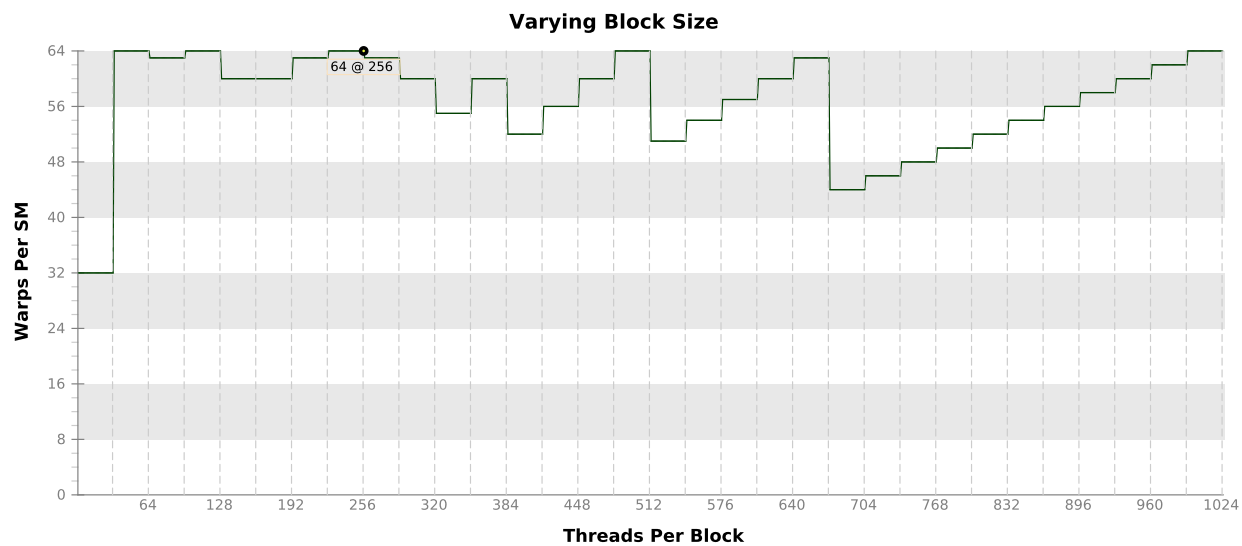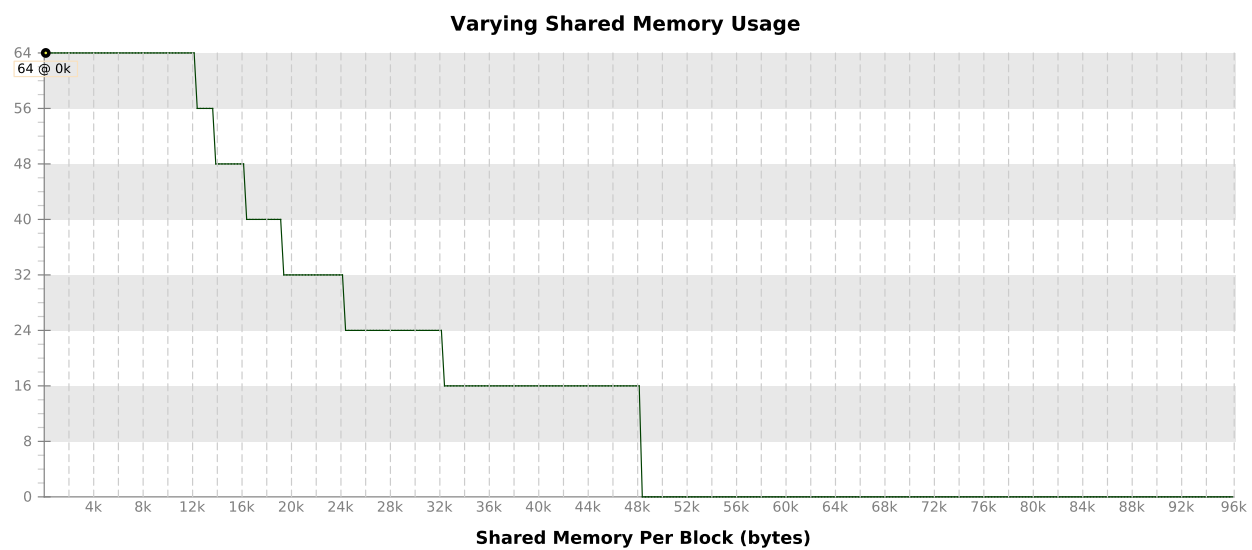
3. Compute Resources

GPU compute resources limit the performance of a kernel when those resources are insufficient or poorly utilized. Compute resources are used most efficiently when all threads in a warp have the same branching and predication behavior. The results below indicate that a significant fraction of the available compute performance is being wasted because branch and predication behavior is differing for threads within a warp.

3.1. Kernel Profile - Instruction Execution

The Kernel Profile - Instruction Execution shows the execution count, inactive threads, and predicated threads for each source and assembly line of the kernel. Using this information you can pinpoint portions of your kernel that are making inefficient use of compute resource due to divergence and predication.

Examine portions of the kernel that have high execution counts and inactive or predicated threads to identify optimization opportunities.

Cuda Fuctions :

void cellHistogramsNaive<unsigned char, unsigned char, int=64, int=8, int=8>(unsigned char*, unsigned char*, int, int, int)

Maximum instruction execution count in assembly: 204
Average instruction execution count in assembly: 179
Instructions executed for the kernel: 27936
Thread instructions executed for the kernel: 822144
Non-predicated thread instructions executed for the kernel: 819152
Warp non-predicated execution efficiency of the kernel: 91.6%
Warp execution efficiency of the kernel: 92.0%

Source files :

/home/adas/cuda-workspace/CudaVisionSysDeploy/Release/../src/init/../device/LBPHist/cellHistograms.h

3.2. Divergent Branches

Compute resource are used most efficiently when all threads in a warp have the same branching behavior. When this does not occur the branch is said to be divergent. Divergent branches lower warp execution efficiency which leads to inefficient use of the GPU's compute resources.

Optimization: Each entry below points to a divergent branch within the kernel. For each branch reduce the amount of intra-warp divergence.

/home/adas/cuda-workspace/CudaVisionSysDeploy/Release/../src/init/../device/LBPHist/cellHistograms.h	
Line 40	Divergence = 23.6% [17 divergent executions out of 72 total executions]

3.3. Function Unit Utilization

Different types of instructions are executed on different function units within each SM. Performance can be limited if a function unit is over-used by the instructions executed by the kernel. The following results show that the kernel's performance is not limited by overuse of any function unit.
 Load/Store - Load and store instructions for shared and constant memory.
 Texture - Load and store instructions for local, global, and texture memory.
 Single - Single-precision integer and floating-point arithmetic instructions.
 Double - Double-precision floating-point arithmetic instructions.
 Special - Special arithmetic instructions such as sin, cos, popc, etc.
 Control-Flow - Direct and indirect branches, jumps, and calls.

3.4. Instruction Execution Counts

The following chart shows the mix of instructions executed by the kernel. The instructions are grouped into classes and for each class the chart shows the percentage of thread execution cycles that were devoted to executing instructions in that class. The "Inactive" result shows the thread executions that did not execute any instruction because the thread was predicated or inactive due to divergence.
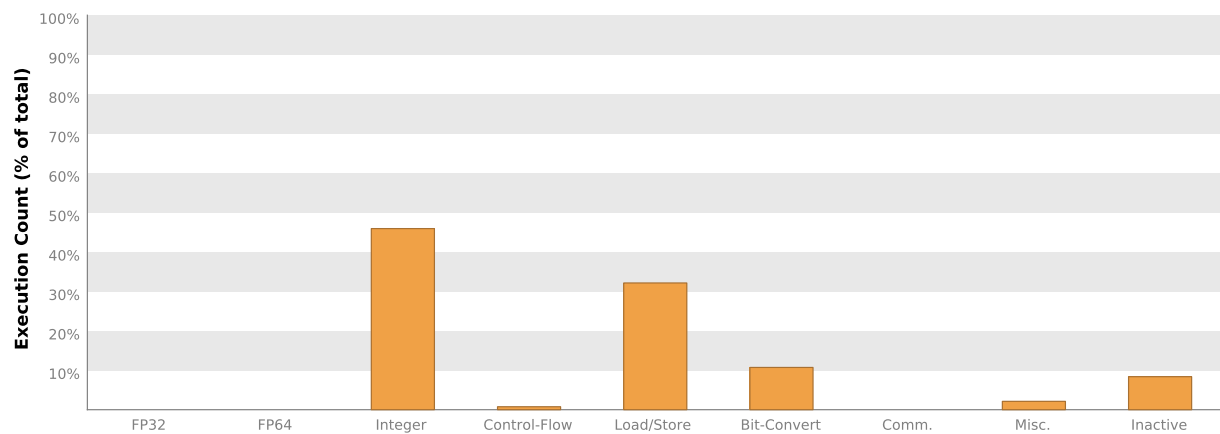
3.5. Floating-Point Operation Counts

The following chart shows the mix of floating-point operations executed by the kernel. The operations are grouped into classes and for each class the chart shows the percentage of thread execution cycles that were devoted to executing operations in that class. The results do not sum to 100% because non-floating-point operations executed by the kernel are not shown in this chart.

4. Memory Bandwidth

Memory bandwidth limits the performance of a kernel when one or more memories in the GPU cannot provide data at the rate requested by the kernel. The results below indicate that the kernel is limited by the bandwidth available to the L2 cache.

4.1. Global Memory Alignment and Access Pattern

Memory bandwidth is used most efficiently when each global memory load and store has proper alignment and access pattern.

Optimization: Each entry below points to a global load or store within the kernel with an inefficient alignment or access pattern. For each load or store improve the alignment and access pattern of the memory access.

/home/adas/cuda-workspace/CudaVisionSysDeploy/Release/../src/init/../device/LBPHist/cellHistograms.h

Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total

	executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total

	executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 7.3, Ideal Transactions/Access = 1 [1496 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Load L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]
Line 48	Global Store L2 Transactions/Access = 29.3, Ideal Transactions/Access = 1 [5984 L2 transactions for 204 total executions]

4.2. Memory Bandwidth And Utilization

The following table shows the memory bandwidth used by this kernel for the various types of memory on the device. The table also shows the utilization of each memory type relative to the maximum throughput supported by the memory.

Transactions	Bandwidth	Utilization	
Shared Memory			
Shared Loads	0	0 B/s	
Shared Stores	0	0 B/s	
Shared Total	0	0 B/s	Idle Low Medium High Max
L2 Cache			
Reads	119847	100.165 GB/s	
Writes	95750	80.025 GB/s	
Total	215597	180.19 GB/s	Idle Low Medium High Max
Unified Cache			
Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Global Loads	119680	100.025 GB/s	
Global Stores	95744	80.02 GB/s	
Texture Reads	26112	21.824 GB/s	
Unified Total	241536	201.869 GB/s	Idle Low Medium High Max
Device Memory			
Reads	40	33.431 MB/s	
Writes	1670	1.396 GB/s	
Total	1710	1.429 GB/s	Idle Low Medium High Max
System Memory			

[PCIe configuration: Gen2 x16, 5 Gbit/s]

Transactions	Bandwidth	Utilization	
Reads	0	0 B/s	Idle Low Medium High Max
Writes	5	4.179 MB/s	Idle Low Medium High Max