

Analysis Report

void computeHOGlocal<float, float, float, int=8, int=8, int=16, int=16, int=64>(float*, float*, float*, float*, int, int, int, int)

Duration	21.696 ms (21,696,147 ns)
Grid Size	[243,1,1]
Block Size	[256,1,1]
Registers/Thread	72
Shared Memory/Block	0 B
Shared Memory Requested	96 KiB
Shared Memory Executed	96 KiB
Shared Memory Bank Size	4 B

[0] GeForce GTX 960


GPU UUID	GPU-0db32734-f94e-48a7-8b5d-4604317dc554
Compute Capability	5.2
Max. Threads per Block	1024
Max. Shared Memory per Block	48 KiB
Max. Registers per Block	65536
Max. Grid Dimensions	[2147483647, 65535, 65535]
Max. Block Dimensions	[1024, 1024, 64]
Max. Warps per Multiprocessor	64
Max. Blocks per Multiprocessor	32
Single Precision FLOP/s	2.644 TeraFLOP/s
Double Precision FLOP/s	82.624 GigaFLOP/s
Number of Multiprocessors	8
Multiprocessor Clock Rate	1.291 GHz
Concurrent Kernel	true
Max IPC	6
Threads per Warp	32
Global Memory Bandwidth	112.16 GB/s
Global Memory Size	4 GiB
Constant Memory Size	64 KiB
L2 Cache Size	1 MiB
Memcpy Engines	2
PCIe Generation	2
PCIe Link Rate	5 Gbit/s
PCIe Link Width	16

1. Compute, Bandwidth, or Latency Bound

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results below indicate that the performance of kernel "void computeHOGlocal<float,..." is most likely limited by memory bandwidth. You should first examine the information in the "Memory Bandwidth" section to determine how it is limiting performance.

1.1. Kernel Performance Is Bound By Memory Bandwidth

For device "GeForce GTX 960" the kernel's compute utilization is significantly lower than its memory utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by the memory system. For this kernel the limiting factor in the memory system is the bandwidth of the L2 Cache memory.



2. Memory Bandwidth

Memory bandwidth limits the performance of a kernel when one or more memories in the GPU cannot provide data at the rate requested by the kernel. The results below indicate that the kernel is limited by the bandwidth available to the L2 cache.

2.1. Global Memory Alignment and Access Pattern

Memory bandwidth is used most efficiently when each global memory load and store has proper alignment and access pattern.

Optimization: Each entry below points to a global load or store within the kernel with an inefficient alignment or access pattern. For each load or store improve the alignment and access pattern of the memory access.

`/home/adas/cuda-workspace/CudaVisionSysDeploy/Release/./src/init/./device/HOG/HOGdescriptor.cuh`

Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]
Line 221	Global Load L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [995312 L2 transactions for 31104 total executions]

/home/adas/cuda-workspace/CudaVisionSysDeploy/Release/./src/init/./device/HOG/HOGdescriptor.cuh

[illegible]

```
/home/adas/cuda-workspace/CudaVisionSysDeploy/Release/./src/init/./device/HOG/HOGdescriptor.cuh
```

[illegible]

/home/adas/cuda-workspace/CudaVisionSysDeploy/Release/./src/init/./device/HOG/HOGdescriptor.cuh

	executions]
Line 230	Global Store L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [62207 L2 transactions for 1944 total executions]
Line 230	Global Store L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [62207 L2 transactions for 1944 total executions]
Line 230	Global Store L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [62207 L2 transactions for 1944 total executions]
Line 230	Global Store L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [62207 L2 transactions for 1944 total executions]
Line 230	Global Store L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [62207 L2 transactions for 1944 total executions]
Line 230	Global Store L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [62207 L2 transactions for 1944 total executions]
Line 230	Global Store L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [62207 L2 transactions for 1944 total executions]
Line 230	Global Store L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [62207 L2 transactions for 1944 total executions]
Line 230	Global Store L2 Transactions/Access = 32, Ideal Transactions/Access = 4 [62207 L2 transactions for 1944 total executions]

2.2. High Local Memory Overhead

Local memory loads and stores account for 75% of total memory traffic. High local memory traffic typically indicates excessive register spilling.

Optimization: Use the -maxrregcount flag or the __launch_bounds__ qualifier to increase the number of registers available to nvcc when compiling the kernel.

2.3. GPU Utilization Is Limited By Memory Bandwidth

The following table shows the memory bandwidth used by this kernel for the various types of memory on the device. The table also shows the utilization of each memory type relative to the maximum throughput supported by the memory. The results show that the kernel's performance is potentially limited by the bandwidth available from one or more of the memories on the device.

Optimization: Try the following optimizations for the memory with high bandwidth utilization.







Shared Memory - If possible use 64-bit accesses to shared memory and 8-byte bank mode to achieved 2x throughput.

L2 Cache - Align and block kernel data to maximize L2 cache efficiency.

Unified Cache - Reallocate texture data to shared or global memory. Resolve alignment and access pattern issues for global loads and stores.

Device Memory - Resolve alignment and access pattern issues for global loads and stores.

System Memory (via PCIe) - Make sure performance critical data is placed in device or shared memory.

Transactions	Bandwidth	Utilization	
Shared Memory			
Shared Loads	0	0 B/s	
Shared Stores	0	0 B/s	
Shared Total	0	0 B/s	
L2 Cache			
Reads	78371465	116.592 GB/s	
Writes	59891931	89.1 GB/s	
Total	138263396	205.693 GB/s	
Unified Cache			
Local Loads	57799900	85.988 GB/s	
Local Stores	57652454	85.769 GB/s	
Global Loads	33840640	48.123 GB/s	
Global Stores	2239452	3.332 GB/s	
Texture Reads	15209856	22.628 GB/s	
Unified Total	166742302	245.839 GB/s	
Device Memory			
Reads	16171696	24.058 GB/s	
Writes	19431782	28.908 GB/s	
Total	35603478	52.967 GB/s	
System Memory			
[PCIe configuration: Gen2 x16, 5 Gbit/s]			
Reads	0	0 B/s	
Writes	5	7.438 kB/s	

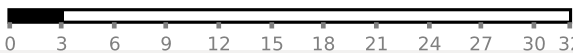
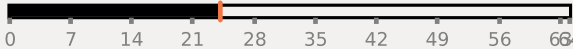

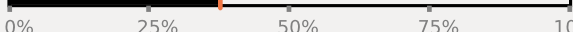

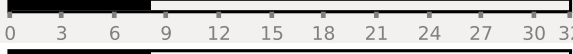
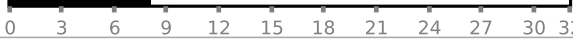


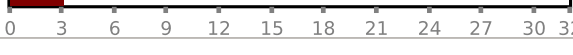

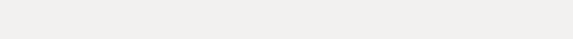
3. Instruction and Memory Latency

Instruction and memory latency limit the performance of a kernel when the GPU does not have enough work to keep busy. The performance of latency-limited kernels can often be improved by increasing occupancy. Occupancy is a measure of how many warps the kernel has active on the GPU, relative to the maximum number of warps supported by the GPU. Theoretical occupancy provides an upper bound while achieved occupancy indicates the kernel's actual occupancy. The results below indicate that occupancy can be improved by reducing the number of registers used by the kernel.

3.1. GPU Utilization Is Limited By Register Usage

The kernel uses 72 registers for each thread (18432 registers for each block). This register usage is likely preventing the kernel from fully utilizing the GPU. Device "GeForce GTX 960" provides up to 65536 registers for each block. Because the kernel uses 18432 registers for each block each SM is limited to simultaneously executing 3 blocks (24 warps). Chart "Varying Register Count" below shows how changing register usage will change the number of blocks that can execute on each SM.


Optimization: Use the `-maxrregcount` flag or the `__launch_bounds__` qualifier to decrease the number of registers used by each thread. This will increase the number of blocks that can execute on each SM. On devices with Compute Capability 5.2 turning global cache off can increase the occupancy limited by register usage.

Variable	Achieved	Theoretical	Device Limit	Grid Size: [243,1,1] (243 blocks) Block Size: [256,1,1] (256 threads)
Occupancy Per SM				
Active Blocks		3	32	
Active Warps	23.75	24	64	
Active Threads		768	2048	
Occupancy	37.1%	37.5%	100%	
Warps				
Threads/Block		256	1024	
Warps/Block		8	32	
Block Limit		8	32	
Registers				
Registers/Thread		72	255	
Registers/Block		18432	65536	
Block Limit		3	32	
Shared Memory				
Shared Memory/Block		0	98304	
Block Limit			32	


3.2. Occupancy Charts


The following charts show how varying different components of the kernel will impact theoretical occupancy.

Varying Block Size



Varying Register Count





4. Compute Resources

GPU compute resources limit the performance of a kernel when those resources are insufficient or poorly utilized.

4.1. Function Unit Utilization

Different types of instructions are executed on different function units within each SM. Performance can be limited if a function unit is over-used by the instructions executed by the kernel. The following results show that the kernel's performance is not limited by overuse of any function unit.

Load/Store - Load and store instructions for shared and constant memory.


Texture - Load and store instructions for local, global, and texture memory.

Single - Single-precision integer and floating-point arithmetic instructions.

Double - Double-precision floating-point arithmetic instructions.


Special - Special arithmetic instructions such as sin, cos, popc, etc.

Control-Flow - Direct and indirect branches, jumps, and calls.



4.2. Instruction Execution Counts

The following chart shows the mix of instructions executed by the kernel. The instructions are grouped into classes and for each class the chart shows the percentage of thread execution cycles that were devoted to executing instructions in that class. The "Inactive" result shows the thread executions that did not execute any instruction because the thread was predicated or inactive due to divergence.



4.3. Floating-Point Operation Counts

The following chart shows the mix of floating-point operations executed by the kernel. The operations are grouped into classes and for each class the chart shows the percentage of thread execution cycles that were devoted to executing operations in that class. The results do not sum to 100% because non-floating-point operations executed by the kernel are not shown in this chart.

